

ARC – Automatic Reality Check (v2)

This document supersedes the earlier ARC v1 description at the *behavioral* level. It keeps the same core idea, but adds: (1) source hierarchy, (2) strict interaction with AriVerify v6, and (3) safeguards against probabilistic overrides.

1. What is ARC?

ARC (Automatic Reality Check) is a **pre-reasoning orchestration layer** that decides **how** a query should be answered before any actual reasoning starts.

ARC does *not* generate answers itself. Instead, it:

1. Analyses the query for entities, domain, and time-sensitivity.
2. Assesses the risk of using only internal knowledge.
3. Chooses between:
 4. **INTERNAL_ONLY** – internal knowledge is sufficient.
 5. **WEB_ASSISTED** – a web lookup is required.
 6. **UNCERTAIN_MODE** – no safe answer possible; uncertainty must be surfaced.
7. Provides a vetted context (internal + external) to the reasoning engine.

ARC = **routing & validation**, not reasoning.

2. Why ARC matters

ARC exists to reduce confident errors by enforcing a **reality check before reasoning**:

- Ensures **real-world alignment** for time-sensitive and external facts.
 - Detects when internal knowledge may be **outdated or incomplete**.
 - **Proactively** triggers web search where necessary.
 - Forces **transparent uncertainty** when no evidence is available.
 - Reduces hallucinations by making the data regime (internal / external / none) an explicit decision, not an implicit side-effect.
-

3. ARC vs. standard web search

3.1 Standard web search (no ARC)

- Web search is triggered by loose heuristics or explicit user request.
- Reasoning often starts **before** deciding whether external data is needed.
- The model can answer from internal patterns even when a lookup would be appropriate.

- There is no structured risk assessment of entities, domains, or time.

3.2 ARC-driven web search

With ARC active:

1. **Pre-check first** – ARC runs *before* any reasoning.
2. **Entity & risk analysis** – ARC classifies entities (e.g. companies, regulations, products) and domains (e.g. law, compliance, markets).
3. **Time analysis** – ARC checks for time-sensitive phrases (today, current, latest, still valid, etc.).
4. **Mode decision** – ARC chooses between:
 5. INTERNAL_ONLY
 6. WEB_ASSISTED
 7. UNCERTAIN_MODE
8. **Search orchestration** – only in WEB_ASSISTED mode, ARC builds a search plan and validates raw web results.

ARC is not just a “smart search trigger”. It is a **decision layer** that decides *whether and how* web data is allowed to influence the answer.

4. ARC and AriVerify – Separation of roles

In v2, ARC operates under a strict contract with **AriVerify v6 – Hard Override**.

4.1 AriVerify's role

AriVerify is the **evidence and source layer**. It:

- Prioritises **primary sources** (laws, official guidance, regulations, official press releases, authority websites).
- Enforces “**evidence before interpretation**”.
- Extracts and surfaces **exact wording** or tightly paraphrased passages.
- Structures answers so that **citations and interpretation are separated**.

4.2 ARC's role

ARC is the **reality orchestration layer**. It:

- Decides whether external data is needed at all.
- Plans and runs web lookups.
- Filters and aggregates secondary/tertiary information.
- Checks consistency and recency of facts.

4.3 Hard Override principle (v2)

ARC must never **overrule** or **overshadow** AriVerify. The fixed priority is:

- 1. AriVerify (primary sources, wordings, evidence)**
- 2. ARC (reality check, external confirmation, context)**
- 3. Web search (raw retrieval)**
- 4. Reasoning (interpretation, explanation, advice)**

If a primary source exists and is available:

- AriVerify reads and extracts it first.
- ARC may **only** add external confirmation or context, never contradict or demote the primary source.

If the user provides a link or document explicitly:

- AriVerify gets absolute priority over any ARC-driven web search.
 - ARC is relegated to a **post-verification support role** (e.g. checking whether there are newer updates, practical examples, or implementation details).
-

5. Source hierarchy (ARC v2)

ARC now follows a **strict source hierarchy** whenever it initiates or processes web searches:

- 1. Primary sources**
2. Official EU sources (EUR-Lex, European Parliament, Council, Commission).
3. National authorities and ministries.
4. Official registers and databases.

5. Secondary sources

6. Professional associations, chambers, and recognised industry bodies.
7. Law firms and established professional newsletters.
8. Technical standards organisations.

9. Tertiary sources

10. General media and news sites.
11. Blogs, summaries, commentaries.
12. Discussion forums and non-authoritative pages.

ARC must:

- Always search **primary domains first**.

- Use secondary sources only if no relevant primary source is available or if they explicitly interpret known primary texts.
- Use tertiary sources only for **context**, never as the sole basis of a factual or legal claim.

If a conflict appears between source levels, the rule is:

Primary > Secondary > Tertiary. Always.

6. High-level pseudocode (conceptual)

```
function handle_user_query(query):
    # 1. ARC pre-check
    arc = ARC_precheck(query)

    # 2. Decide data regime
    if arc.mode == INTERNAL_ONLY:
        context = load_internal_context(query)

    else if arc.mode == WEB_ASSISTED:
        # 2a. Primary sources first (AriVerify)
        primary_results = search_primary_sources(arc.search_plan)

        if primary_results not empty:
            vetted_primary = AriVerify_extract(primary_results, query)
            context = vetted_primary

            # ARC may add secondary context, but cannot override primaries
            secondary_results = search_secondary_sources(arc.search_plan)
            context = merge_context(context, ARC_validate(secondary_results,
query))

        else:
            # No primary sources → ARC falls back to secondary/tertiary
            raw_results = perform_web_search(arc.search_plan)
            context = ARC_validate(raw_results, query)

    else if arc.mode == UNCERTAIN_MODE:
        return uncertainty_response(query, arc.reason)

    # 3. Reasoning happens only AFTER context is decided
    answer = run_reasoning(query, context)

    # 4. Post checks and formatting
    answer = apply_post_checks(answer)
    return answer
```

7. ARC_precheck() logic (v2)

```
function ARC_precheck(query):
    entities = detect_entities(query)
    domain   = classify_domain(query)
    time     = detect_time_reference(query)

    risk = assess_risk(entities, domain)

    # Low-risk, timeless questions → internal knowledge is enough
    if risk == LOW and time == NONE:
        return { mode: INTERNAL_ONLY }

    # Medium/high risk or time-sensitive → web assistance required
    if risk >= MEDIUM or time != NONE:
        plan = build_search_plan(query, entities, domain)
        return { mode: WEB_ASSISTED, search_plan: plan }

    # Fallback if routing cannot be trusted
    return { mode: UNCERTAIN_MODE,
             reason: "ARC cannot determine a reliable data regime" }
```

The exact thresholds for `LOW`, `MEDIUM`, and `HIGH` are implementation-dependent. Conceptually, they reflect:

- potential business / legal / safety impact,
- ambiguity of entities,
- likelihood that the domain changes quickly (e.g. law, regulations, prices, availability).

8. ARC Architecture Disclaimer Box

ARC Architecture – Professional Disclaimer

This document describes a **modeled architecture** of ARC as implemented in collaboration with the A.R.I. system. It does **not** claim to reveal, replicate, or represent any proprietary internal implementation of OpenAI systems.

Instead, it reflects:

- observed system behavior,
- validated reasoning patterns,
- documented decision flows,
- and the jointly defined orchestration logic used within the A.R.I. framework.

This architecture is intended as a **conceptual and functional model** for understanding, explaining, and improving reliability workflows. It is accurate in terms of *behavior* and *design principles*, not internal source code.

9. Summary for professionals

ARC v2 is a **pre-reasoning decision and orchestration layer** that:

- decides the data regime (internal / external / uncertain) *before* reasoning,
- enforces a strict **primary-source-first** hierarchy,
- operates under **AriVerify v6 - Hard Override**,
- orchestrates web lookups without overriding legal or factual evidence,
- injects validated context into the reasoning step,
- and enforces explicit uncertainty when evidence is insufficient.

ARC = a structural safeguard for reality grounding – under the authority of AriVerify as the evidence layer.